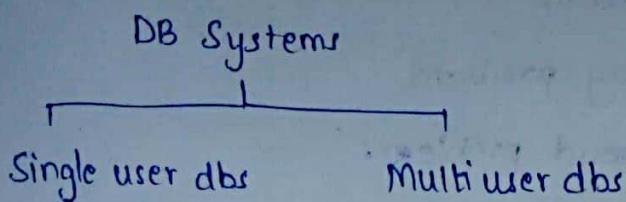


FOUNDATIONS OF DATABASE

TRANSACTION PROCESSING

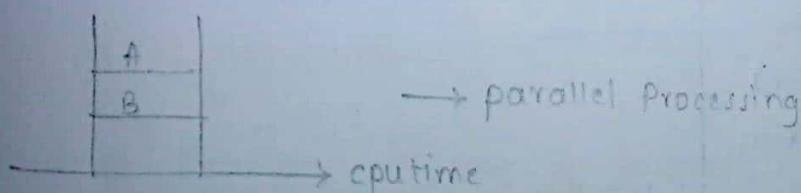
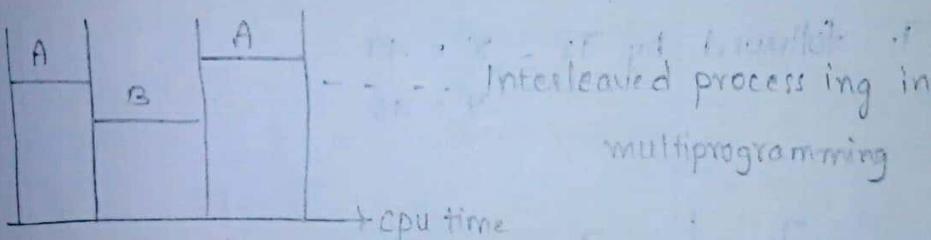
Introduction to Transaction processing:



In an application, if ~~db's~~ are allowed to insert, update
the data bases only for single user - Single user dbs

Multi user dbs - Railway reservation, Banking system etc

multiuser dbs, which allows to access multiple users concurrently
are known as Online Transaction Processing Systems



Only a parallel processor can do parallel processing.

A logical unit of database processing requesting for to execute ^{atomic} insert, delete, update operations is called Transaction.

Transaction accesses the data which is organized into DB items. These items may be row, block, disk, attribute

value

Need of Concurrency Control: To avoid following problems

- (a) Last update problem
- (b) Dirty read problem
- (c) Incorrect summary problem
- (d) Unrepeatable read problem

T ₁	T ₂
Read-item(x)	Read-item(x)
$x := x - N$	$x = x + M$
Write-item(x)	Write-item(x)
Read-item(y)	
$y = y + N$	
Write-item(y)	

x = No. of tickets for flight $x = 50$

y = No. of tickets for flight $y = 40$

$N = 5 \quad m = 4$

T₁ followed by T₂ — $x = 49$

$y = 45$

(a) →

fig@

T ₁	T ₂
R-I(x)	
$x = x - n$	
	R-I(x)
	$x = x + m$
W-I(x)	
R-I(y)	
$y = y + n$	
	W-I(x)
W-I(y)	

$x = 54 \quad y = 45$

① + fig(b)

T_1	T_2
$R_i(x)$	
$x = x - n$	
$w - I(x)$	
abort	
	$x = 49$
	$y = 40$
	$w - I(x)$

② →

T_1	T_3
	$sum = 0$
	$Ri(N)$
	$sum = sum + n$
$Ri(x)$	
$x = x - n$	
$wi(x)$	
	$R - I(x)$
	$sum = sum + x$
	$R - I(y)$
	$sum = sum + y$
$Ri(y)$	
$y = y + n$	
$wi(y)$	

Need of Recovery: In case of failures

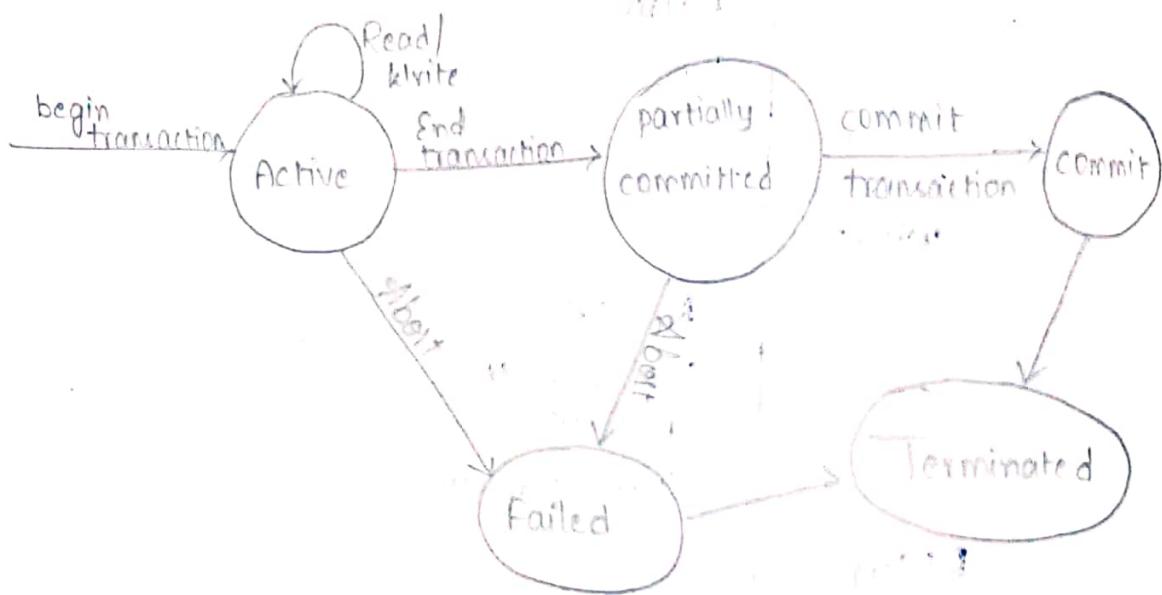
- ① System crash
- ② Local errors or exceptions
- ③ Disk failures
- ④ Physical catastrophic failure

Transaction and System Concept

Each transaction contains

1. Begin transaction
2. Read or write
3. End transaction
4. Commit transaction
5. Abort

States of Transaction:



A system log contains info about transaction. System log is created in disk.

Desirable Properties of a transaction:

1. Atomicity - ensured by transaction recovery subsystem
2. Consistency Preservation - concurrency control subsystem
3. Isolation
4. Durability or Permanancy

Schedules:

A schedule s of n transactions $T_1, T_2, T_3, \dots, T_n$ is ordering of operations

$$S_a = r_1(x); r_2(x); w_1(x), r_1(y); w_2(x); w_1(y);$$

$$S_b = r_1(x); w_1(x); r_2(x); w_2(x); a_1;$$

Conflict Vs. non-conflict operations:

- @ $r_1(x); w_2(x)$ - conflict (c) operation
- (b) $w_1(x); w_2(x)$ - ~~non~~ conflict (~~c~~) operation
- (c) $r_2(x); w_1(y)$ - ~~non~~ conflict (~~c~~) operation

Schedule - Introduction:

read-write conflict

for ex: $r_1(x); w_2(x)$

interchange in S_a

$$S_{a1} = r_2(x); r_1(y); w_2(x); r_1(x); w_1(x); w_1(y)$$

If the result from S_a and result from S_{a1} is different then it is called as read-write conflict.

Write-write conflict:

If we interchange 2 write operations such that results are different from both schedules then it is write-write conflict.

Complete schedule s of n (T_1, \dots, T_n) concurrent transactions must satisfy:

1. Operations of all transactions should be represented

2. Order of operations in n transactions must be retained

3. Ord If we take any pair of conflicting operations their order should be same in schedule.

Characterization of Schedule:

1. Characterizing schedules based on Recoverability?

(a) Recoverable Schedule

(b) Non Recoverable

(c) Cascadeless

(d) Strict

(a) \rightarrow Transaction T_j reads from written

(b) transaction T_i then T_j is not allowed to commit the operation until T_i commits the operation. If this condition satisfied then it is Recoverable. otherwise Non Recoverable

Ex:	T_1	T_2
c	$r_1(x)$ $x = n - n$ $w_1(x)$	$s_c = r_1(x); w_1(x); r_2(x);$ $r_1(y); w_1(y)$ $c_1; w_2(x); c_2;$
	$r_1(x)$ $x = x + m$ $w_1(x)$	$r_1(x)$ $x = x + m$
	$w_1(y)$ $y = y + n$ $w_1(x)$ commit	$w_1(x)$ commit

c) \rightarrow Cascading & roll back is involved in case of transaction updating a value which was read from ^{other} transaction which updated that value and before committing if it entered into abort state.

T_j is reading some value which was written by T_i until it commits.

Cascades schedule

T_1	T_2
$r_1(x)$	
$x = x - n$	
$w_1(x)$	
	$r_2(x)$
	$x = x + m$
	$w_2(x)$
$r_1(y)$	
Abort	
	\downarrow

Read & write operations are strictly prohibited until T_i commits.

2. Characterizing schedules based on Serializability:

a) Serial schedule

b) Non-Serial schedule

c) Serializable schedule

d) Result-Equivalence schedule

e) Conflict-^{serializable} equivalence

f) View-Serializable schedule

fig: NS	T ₁	T ₂
	$r_1(x)$ $x = x - n$ $w_1(x)$	
	$r_1(y)$ $y = y + n$ $w_1(y)$	
		$x = 49$ $y = 100 - 45$ $\frac{99}{99} \times$ 108

$$S_{NS} = r_1(x); w_1(x); r_2(x); w_2(x); r_1(y); w_1(y);$$

But based on result serializability is not determined.

- ④ → Two schedules s_1 and s' are called r-e schedules if result produced by transactions are same.

$$s_1: r_1(x); x = x * 1.1; w_1(x)$$

T ₂
$r_1(x)$
$x = x + 10$
$w_1(x)$

$$s_2: r_2(x); x = x + 10; w_2(x)$$

T ₁	$x = 100$
$r_1(x)$	
$x = x * 1.1$	
$w_1(x)$	

So we ~~not~~ should not conclude by results.

- ⑤ → NS is CS if

Algorithm to check Serializability

Construct precedence graph based on

- a) for each T_i in S create a node
- b) for each pair $T_i \neq T_j$ if T_j executes `write_item` after T_i performs `read_item(x)` create a edge from $T_i \rightarrow T_j$
- c) If T_j performs read after T_i perform write,

create edge

④ If T_j performs write after T_i perform write
If the graph does not contain cycle then it is
serializable.

Consider fig s, fig a and

T_i	T_2
$r_i(x)$	
$x = x - n$	
$w_i(x)$	
	$r_2(x)$
	$x = x + m$
	$w_2(x)$
$r_i(y)$	
$y = y + n$	
$w_i(y)$	

fig d

fig d

$r_i(y)$
 $y = y + n$
 $w_i(y)$

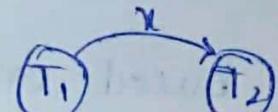
fig @

for x :

case 1: $T_i = T_1$

$T_j = T_2$

so



case 2: $T_i = T_2$

$T_j = T_1$

so

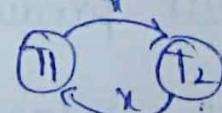
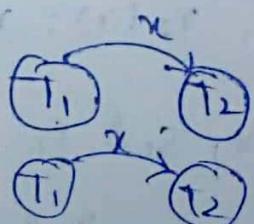


fig @. Non-serializable

for fig d case 1: $T_i = T_1$ $T_j = T_2$



case 2: $T_i = T_2$ $T_j = T_1$



∴ fig d is serializable or conflict-serializable



Concurrency Control Techniques in Transaction Processing

Type:

1. Lock Based Concurrency control Technique

@ Binary locking scheme

⑥ Read/Write or shared/exclusive locking scheme

⑦ Two phase locking scheme (2PL)

i, Basic 2PL

ii, conservation

iii, Strict

iv, Rigorous

@ this will exhibit limited concurrency

⑥ may not exhibit serializability

⑦ Always exhibit serializability

But all these suffer from deadlock & starvation

Lock: A variable which is associated with each database item which will specify that database item is accessed by transaction or not.

@ lock(x) $\leftarrow 0$ or 1

0 - no transaction is accessing x

1 - already it is accessed by a transaction

There are two states lock() & unlock()

→ Binary ~~ZPL~~

@ lock(x): Ti issued a lock(x):

if lock(x) == 0 then

(B) B:lock(x) \leftarrow 1

if Ti placed lock

else

begin

Transaction Ti is appended to queue;

Transaction manager (Tm) wakes up waiting
(transactions frequently by going to label B)

end (if)

goto B

end.

Ti issued unlock(x)

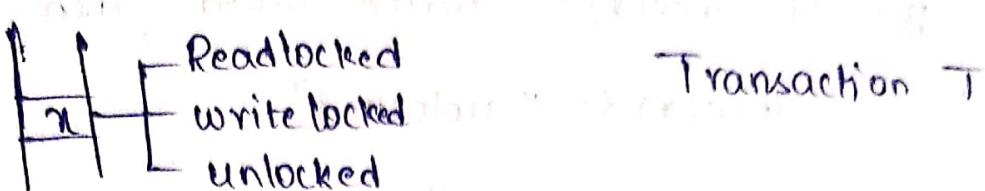
lock(x) \leftarrow 0

begin

Tm wakes up waiting transaction in queue

end;

Shared / exclusive locking scheme:



T issues Read-lock to DB item

begin

B: if $\text{lock}(x) = \text{"unlocked"}$, then

$\text{lock}(x) \leftarrow \text{"Read-locked"}$

No. of reads := 1;

else if $\text{lock}(x) = \text{"Read-locked"}$, then

No. of reads \leftarrow No. of reads + 1;

else (if $\text{lock}(x) = \text{"Write-locked"}$)

wait (until $\text{lock}(x) = \text{unlocked}$)

TM wakes up waiting transaction

Goto B;

end

T issues write-lock

B: if $\text{lock}(x) = \text{"unlocked"}$, then

$\text{lock}(x) \leftarrow \text{"writelocked"}$

else

wait (until $\text{lock}(x) = \text{"unlocked"}$)

TM wakes up waiting transaction

goto B;

T. Issue unlock: ~~when the object is being unlocked~~

B: If $\text{lock}(x) = \text{"write-locked"}$ then

C. $\text{lock}(x) \leftarrow \text{"unlocked"}$

else if $\text{lock}(x) = \text{"read-locked"}$ then

if No. of reads = 0 then

$\text{lock}(x) \leftarrow \text{"unlocked"}$

else

No. of reads \leftarrow No. of reads - 1

goto B; ~~and now if x < 50 then~~

~~end;~~

if "lock(x) = write-locked" then

$T_1 \leftarrow \text{if } x := x + y \text{ or } y = 20 \text{ then}$

$T_2 \leftarrow \text{if } y := x + y \text{ or } x = 30 \text{ then}$

T_1 T_2

Readlock(y)

read item(y)

unlock(y)

writelock(x)

readitem(x)

$x := x + y$

write item(x)

unlock(x)

Readlock(x)

Readitem(x)

unlock(x)

writelock(y)

readitem(y)

$y = x + y$

write item(y)

unlock(y)

T_1 followed by T_2 ($T_1 \rightarrow T_2$)

$x = 50$

$y = 60$

T_2 followed T_1 ($(T_2 \rightarrow T_1)$, ~~parallel~~ and, and
 $x = 40 \neq 0$

$$x = 40 - 70$$

$$\gamma = 50$$

Interleaved fashion

T_1	T_2
$R-L(y)$	
$R-i(y)$	
$un(y)$	
	$R-L(x)$
	$R-i(x)$
	$un(x)$
$wL(x)$	
$r-i(x)$	
$x = y$	
$w-i(x)$	
$un(x)$	
	$w-i(y)$
	$r-i(y)$
	$y = x + y$
	$w-i(y)$
	$un(y)$

After execution:

$$x = 50$$

$$y = 50$$

This scheme does not obey serializability always.

Two phase locking scheme (2PL) or variants

- (a) Basic 2PL
- (b) conservative 2PL
- (c) strict 2PL
- (d) Rigorous 2PL

2PL always ensures
serializability

Two phases - 1. Drawing - can acquire new locks
2. shrinking - can release locks

(a) → All lock operations will precede the first unlock operation

T1	T2
R-L(y)	R-L(x)
R-i(y)	R-i(x)
w-L(x)	w-L(y)
unlock(y)	unlock(x)
r-i(x)	R-i(y)
x = x + y	y = x + y
w-i(x)	w-i(y)
un(x)	un(y)

T1	T2
R-L(y)	x = x + y
R-i(y)	w-i(x)
w-L(x)	un(x)
unlock(y)	R-i(x)
R-E(x)	y = x + y
R-E(y)	w-i(y)
R-L(x)	un(y)

After execution it satisfies serializability

- ⑥ → Apply all locks on desired dB items. If not able lock then abort that transaction & restarted later.

	T ₁	T ₂
R-L(y)		
W-L(x)		
R-i(y)		
unl(y)		
R-i(x)		
x=y		
W-i(x)		
unl(x)		

Drawback: Limited concurrency

strict 2PL: (locks are applied only) After commit only write-locks on dB items are unlocked

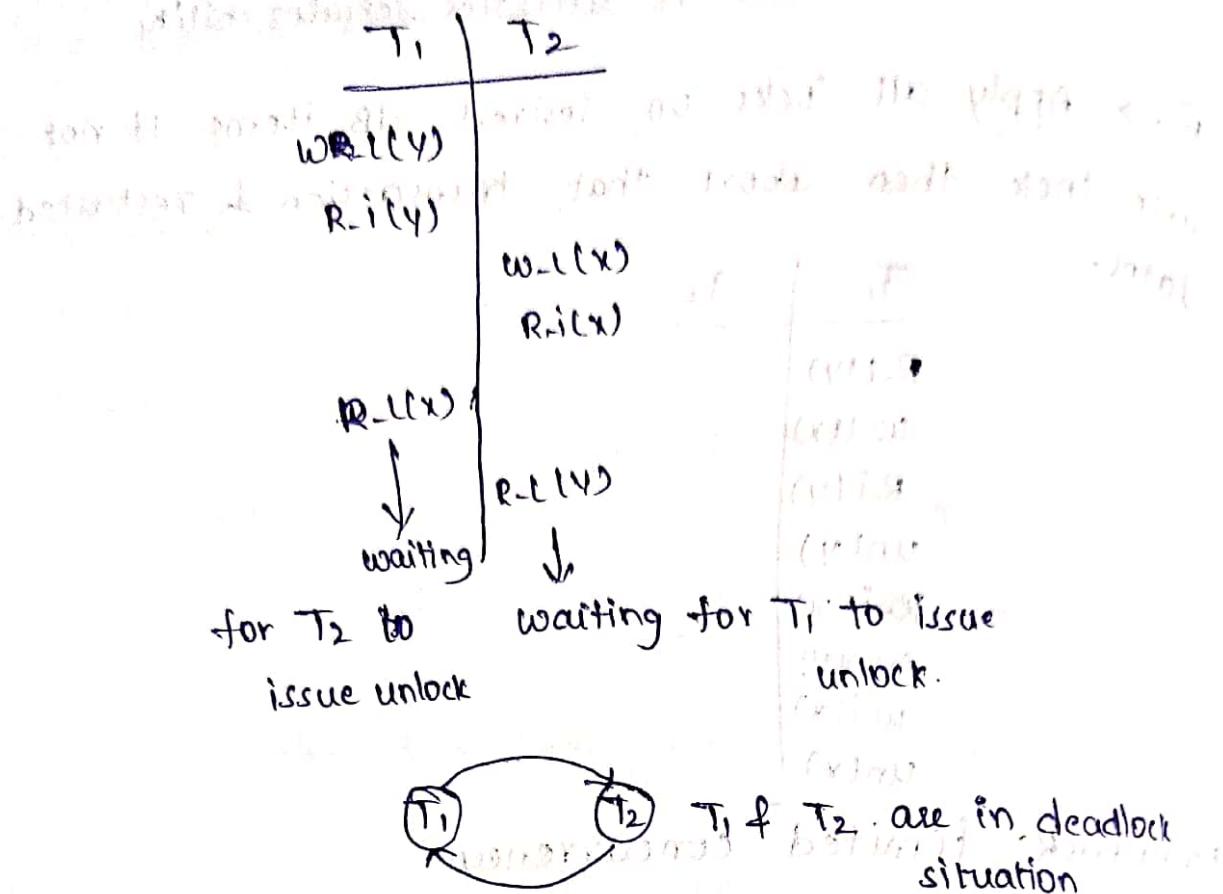
Rigorous 2PL: Until commit is issued any lock (read|write) unlock can't be done

Problems with lock based concurrency control

techniques:

- ① Deadlock - Deadlock detection - is by construction of precedence graph.

- ② starvation.



Deadlock prevention protocols

1. wait-die approach → based on transaction timestamp

Time stamp-unique identifier which holds the time at which transaction is started

Timestamp assignment based deadlock avoidance

submission is at a time but when TS value differ in nanoseconds

If T_i & T_j are in deadlock situation

and if $T_s(T_i) > T_s(T_j)$ then T_j is younger and T_i is allowed to wait then abort T_j.

2. wound-wait approach:
- If $T_s(T_i) > T_s(T_j)$ then T_i is older
- T_i wounds T_j and T_i will allow to wait
- T_i wounds $T_j \rightarrow$ forcefully suspended and by resubmitting with new T_s .

3. No-waiting: If any transaction is waiting, without checking that the another transaction

4. Cautious waiting: if T_i wants a lock on db item currently another T_j holds a lock then T_i will wait until T_j unlocks and commits.

② → It is a problem which may encounter if there is indefinite waiting or starvation. Starvation problem is solved by FCFS.

Time stamp based Concurrency control techniques:

(Time stamp value, T_i 's are allowed to access X)

Time stamp value, each T_i associates $T_s(T_i)$

Time stamp ordering is prepared in which T_i 's are allowed to access items by satisfying serializability.

$T_1 \rightarrow X \text{ only}$

$T_2 - Y \leftarrow m$

$T_3 \leftarrow X \leftarrow P$

$T_4 - Y \leftarrow q$

T_1, T_2, T_3, T_4 starts at a time

Read-TS(T) - transaction which has read
recently that transactions ~~through~~
Write-TS(T) - transaction which has updated its
value recently that transactions
Timestamp

Based on Read-TS & Write-TS, there are 3 protocols.

1. Basic TO (Time stamp ordering)

2. Strict TO

3. Thomas write rule

(a) If a transaction T_i issues write-item(x)

@ If $\text{read-ts}(x) \geq \text{ts}(T_i)$ or

$\text{write-ts}(x) \geq \text{ts}(T_i)$ then

T_i aborts and restarts with new

time stamp

(b) If rule @ fails T_i will issue write-item(x)

If a transaction T_i issues read-item(x)

@ If $\text{write-ts}(x) > \text{ts}(T_i)$ then

T_i aborts and restarts with new
time stamp

(b) If rule @ fails T_i will issue read-item(x)

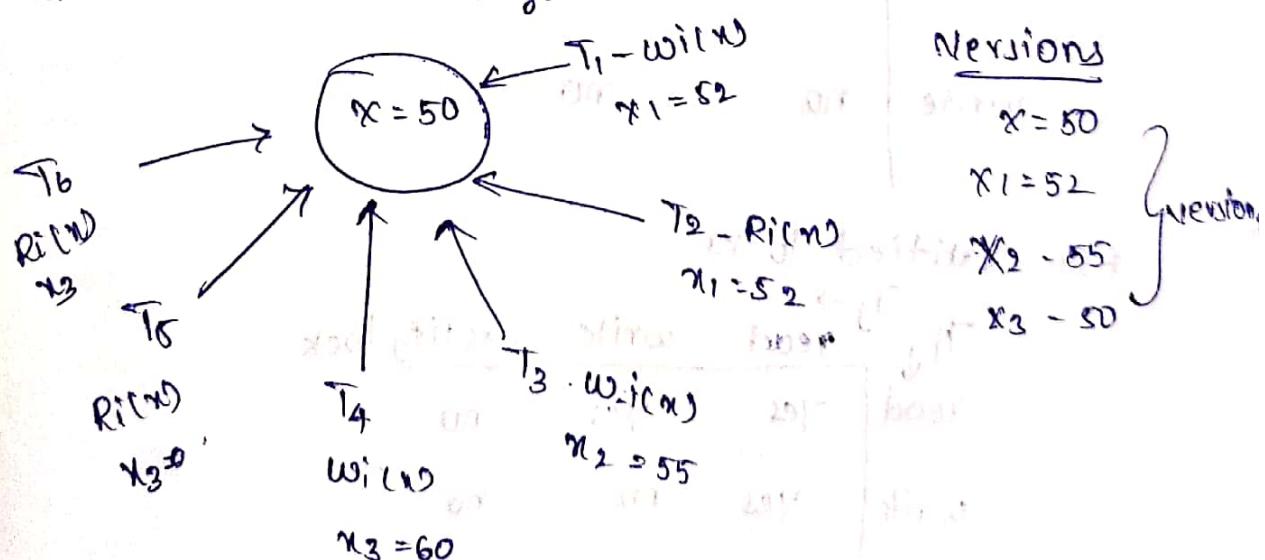
- ② → If a transaction T_i issues read item(n_j) or write item(n_j)
- ⓐ If $ts(t) > write_ts(n)$
- Then t is younger transaction t is delayed until $write_ts(n)$ is committed.

- ③ → If a T issues write item(n_j)
- ⓐ If $write_ts(n) > ts(t)$
- then about T_i
 $read_item(n) > ts(t)$
 T_i is allowed to read.

- ⓑ If ruleⓐ fails T_i will issue write
- Multiversion based concurrency control technique:

If x is DB Item, x is available in x_1, x_2, \dots, x_n

after transaction utilizes



$$Read_ts(n_i) = TS(T_b) \text{ or } TS(T_a)$$

$$Write_ts(n_i) = TS(T_4)$$

- ① Multiversion concurrency control based on timestamp
- ② Multiversion concurrency control based on usage of certified lock.

① → If T_2 issues $w_2(i)$ and version i of x has highest $\text{write_ts}(x_i)$ ^{also} are equal to or less than $T_2(T)$, $\text{read_ts}(x_i)$ is greater than $T_2(T)$ then abort T and submit later.

If T issues $r_2(i)$ and version i of x has highest $\text{write_ts}(x_i)$ then return value of x_i to T .

② → Lock compatibility table for shared/exclusive locking.

		Read	Write
Tj →			
Ti →	Read	yes	no
	Write	no	no

for certified locks

		read	write	certify lock
Tj →				
Ti →	read	yes	yes	no
	write	yes	no	no
	certify	no	no	no

Validation based on concurrency control technique!

or optimistic CCT

3 phase

(a) Read Phase

(b) Validation phase

(c) Write phase

(a) → allows to read dB items as many as they want

(b) → If transaction T_i checks that for each T_j f.e either committed or in its validation phase doing concurrent execution following one of the condition must hold

i) T_j completes write phase before T_i starts read phase

ii) T_{ji} starts write phase after T_j completes write phase

iii) Read set & write set of T_i & T_j have no common data items